

Ray: A Distributed Framework for Emerging AI Applications

Saber Malekmohammadi

November 2021

Computer Science Department




Background - AI applications and system requirements

- Different applications of AI have their **specific computational tasks**
- Based on these tasks, they impose some **system requirements**
- Example: supervised Learning application:

The **stateful** training task

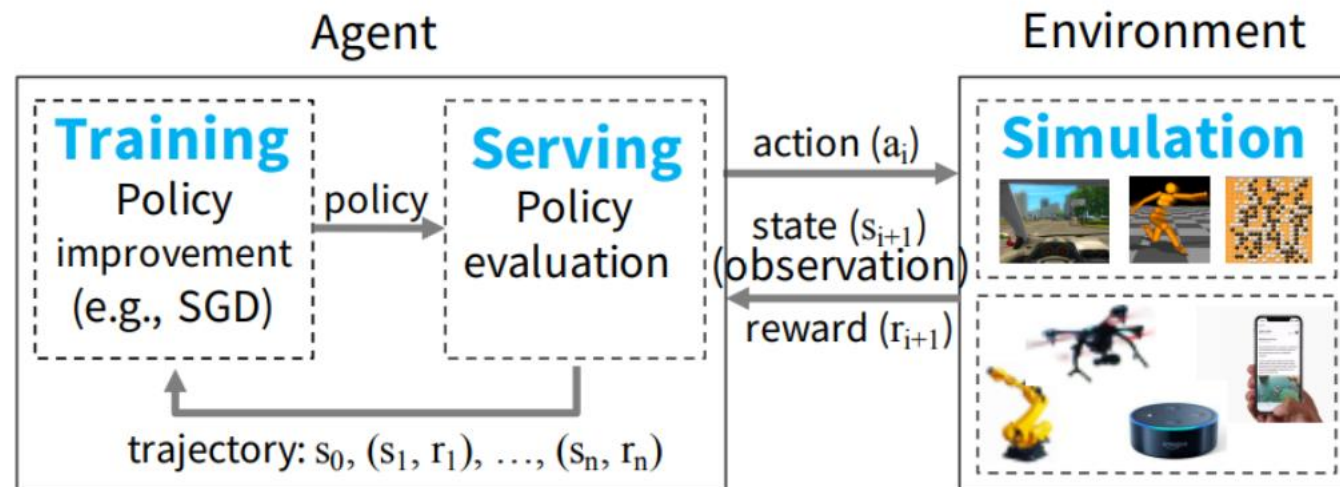
The **stateless** prediction task



Impose system requirements (training stage):
Tensorflow, MXNet and Pytorch

Research Problem

- The scope of AI applications encompasses **more complex applications**
- Three simultaneously required capabilities in Reinforcement Learning (RL):
 - **Distributed training**: **fine-grained computations, heterogeneous computations**
 - **Serving**: **latency-sensitive, fine-grained computations, heterogeneous computations**
 - **Simulations**: **dynamic execution**



Prior Work

- The existing frameworks developed for supervised learning and Big Data workloads can not meet these requirements
- Bulk-synchronous parallel systems such as MapReduce, Apache Spark and Dryad:
Do not support fine-grained action rendering and simulation computations in RL
- Task-parallel systems such as CIEL and DASK:
Do not completely support distributed training and serving
- Distributed deep learning frameworks such as Ten-sorflow and MXNet:
Do not support simulation and serving naturally

Solution

- Considering the system requirements of more complex applications such as RL:
 - Ray, a distributed framework, was proposed for RL applications requirements
 - Provides a general programming model supporting **task-parallel** and **actor-based computations**
 - Supports a range of computations: from **lightweight and stateless** computations (simulations) to **long and stateful** computations (training)
 - Provides **low latency**, **high scalability** and **failure tolerance**
- This enables us to do all the tasks of training, serving and simulation together by a single framework.

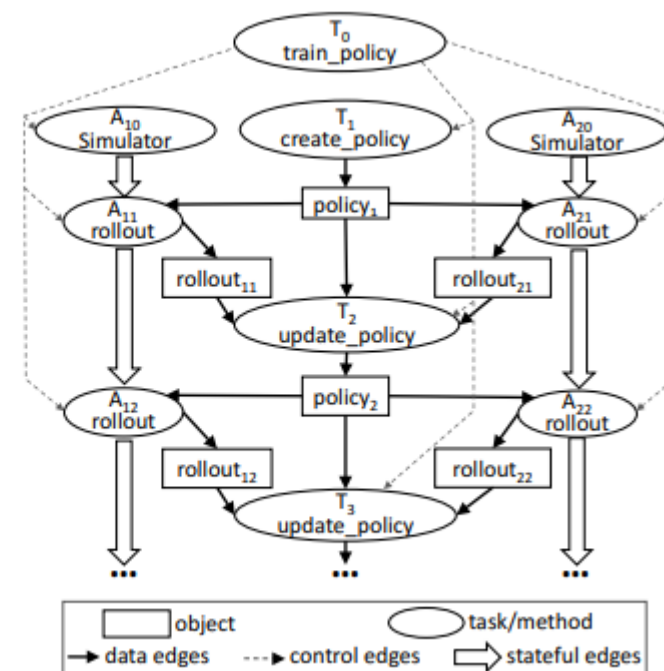
Technical Details – Dynamic Computation Graph

```
@ray.remote
def create_policy():
    # Initialize the policy randomly.
    return policy
```

```
@ray.remote(num_gpus=1)
class Simulator(object):
    def __init__(self):
        # Initialize the environment.
        self.env = Environment()
    def rollout(self, policy, num_steps):
        observations = []
        observation = self.env.current_state()
        for _ in range(num_steps):
            action = compute(policy, observation)
            observation = self.env.step(action)
            observations.append(observation)
        return observations
```

```
@ray.remote(num_gpus=2)
def update_policy(policy, *rollouts):
    # Update the policy.
    return policy
```

```
@ray.remote
def train_policy():
    # Create a policy.
    policy_id = create_policy.remote()
    # Create 10 actors.
    simulators = [Simulator.remote() for _ in range(10)]
    # Do 100 steps of training.
    for _ in range(100):
        # Perform one rollout on each actor.
        rollout_ids = [s.rollout.remote(policy_id) for s in simulators]
        # Update the policy with the rollouts.
        policy_id = update_policy.remote(policy_id,
                                         *rollout_ids)
    return ray.get(policy_id)
```

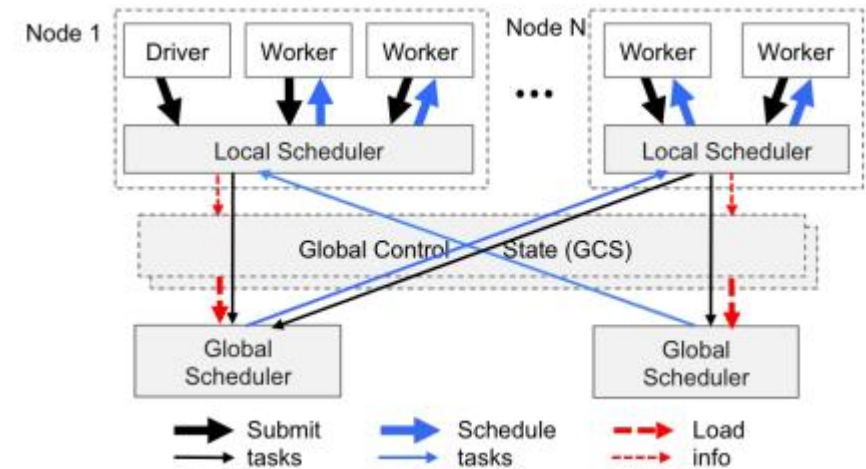
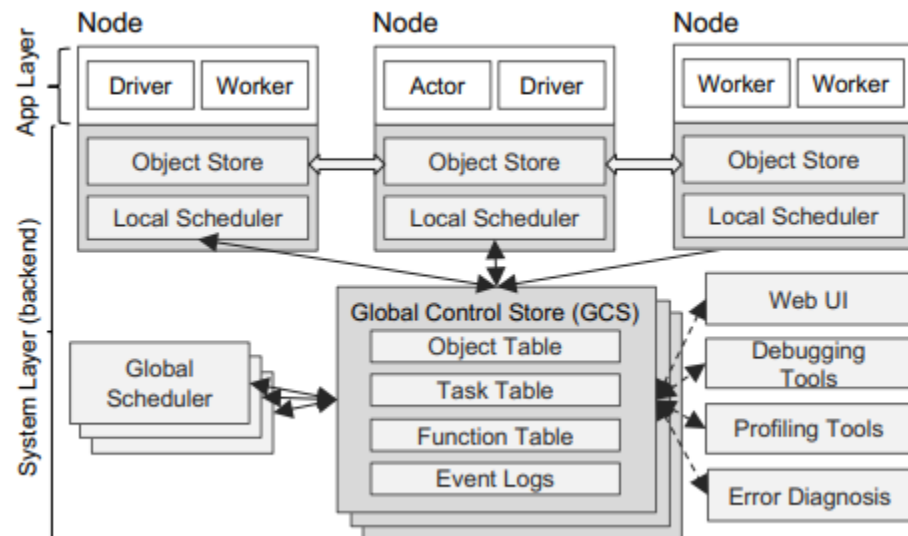


Tasks (stateless)	Actors (stateful)
Fine grained load balancing	Coarse grained load balancing
Support for object locality	Poor locality support
High overhead for small updates	Low overhead for small updates
Efficient failure handling	Overhead from checkpointing

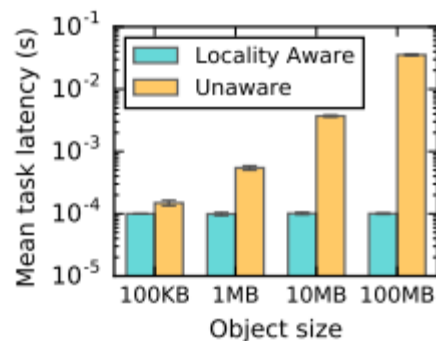
Name	Description
<code>futures = f.remote(args)</code>	Execute function <i>f</i> remotely. <code>f.remote()</code> can take objects or futures as inputs and returns one or more futures. This is non-blocking.
<code>objects = ray.get(futures)</code>	Return the values associated with one or more futures. This is blocking.
<code>ready_futures = ray.wait(futures, k, timeout)</code>	Return the futures whose corresponding tasks have completed as soon as either <i>k</i> have completed or the timeout expires.
<code>actor = Class.remote(args)</code> <code>futures = actor.method.remote(args)</code>	Instantiate class <i>Class</i> as a remote actor, and return a handle to it. Call a method on the remote actor and return one or more futures. Both are non-blocking.

Technical Details – Architecture


- **Locality aware task scheduling**: load balancing and locality-aware scheduling
- **Global control state**: scalability, task/actor failure tolerance
- **Separating task scheduling from task dispatch**: scalability, high throughput of fine-grained tasks, low latency
- **Distributed object store**: unifying actors and tasks on nodes




Performance




(a) Ray locality scheduling




(b) Ray scalability




(a) Task reconstruction



(b) Actor reconstruction



(a) Evolution Strategies



(b) PPO

System	Small Input	Larger Input
Clipper	4400 ± 15 states/sec	290 ± 1.3 states/sec
Ray	6200 ± 21 states/sec	6900 ± 150 states/sec

Strengths and Weaknesses

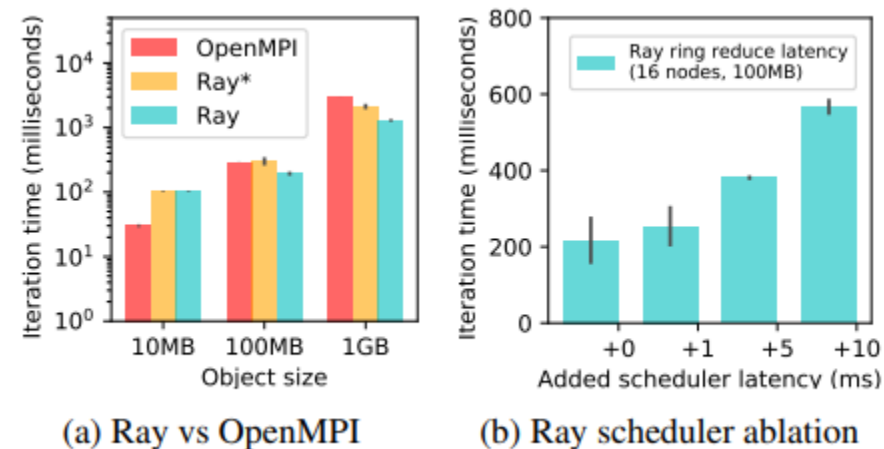
■ Strengths:

- Supporting task-parallel and actor-based computations: supporting training, serving, simulation
- Locality aware task scheduling: load balancing and locality-aware scheduling
- Global control state: scalability, task/actor failure tolerance
- Separating task scheduling from task dispatch: scalability, high task throughput, low latency
- Distributed object store: unifying actors and tasks on nodes
- Scalable architecture: scheduler

■ Weaknesses:

- Higher overhead imposed on distributed primitives:

For small objects, Ray can not outperform dedicated systems like OpenMP



Discussion

- The Comparison with dedicated systems is just based on RL workloads
- No comparison with other workloads
- Object store does not support distributed objects storage: distributed objects like large matrices or trees that do not fit on a single node. Would that impose large latencies?

References

- Joseph E. Gonzalez, et al. **“GraphX: Graph Processing in a Distributed Dataflow Framework”**. In: OSDI. 2014.
- Derek Gordon Murray, et al. **“Naiad: a timely dataflow system”**. In: Proceedings of the ACM Symposium on Operating Systems Principles (2013).
- Jeffrey Dean, et al. **“MapReduce: simplified data processing on large clusters”**. In: Commun. ACM 51 (2008), pp. 107–113.
- Mart´ın Abadi, et al. **“TensorFlow: A system for large-scale machine learning”**. In: OSDI. 2016.
- Tianqi Chen, et al. **“MXNet: A Flexible and Efficient Machine Learning Library for Heterogeneous Distributed Systems”**. In: ArXiv abs/1512.01274 (2015).
- Adam Paszke, et al. **“Automatic differentiation in PyTorch”**. In: 2017.
- Philipp Moritz, et al. **“Ray: A Distributed Framework for Emerging AI Applications”**. In: OSDI. 2018.
- Derek Gordon Murray, et al. **“CIEL: A Universal Execution Engine for Distributed Data-Flow Computing”**. In: NSDI. 2011.
- Matthew Rocklin. **“Dask: Parallel Computation with Blocked algorithms and Task Scheduling”**. In: 2015.